

Big Data

Hands on Spark


A definition for Big Data

"Big data **exceeds the reach of commonly used hardware** environments **and software** tools to capture, manage, and process it with in a tolerable elapsed time for its user population." - *Teradata Magazine article, 2011*

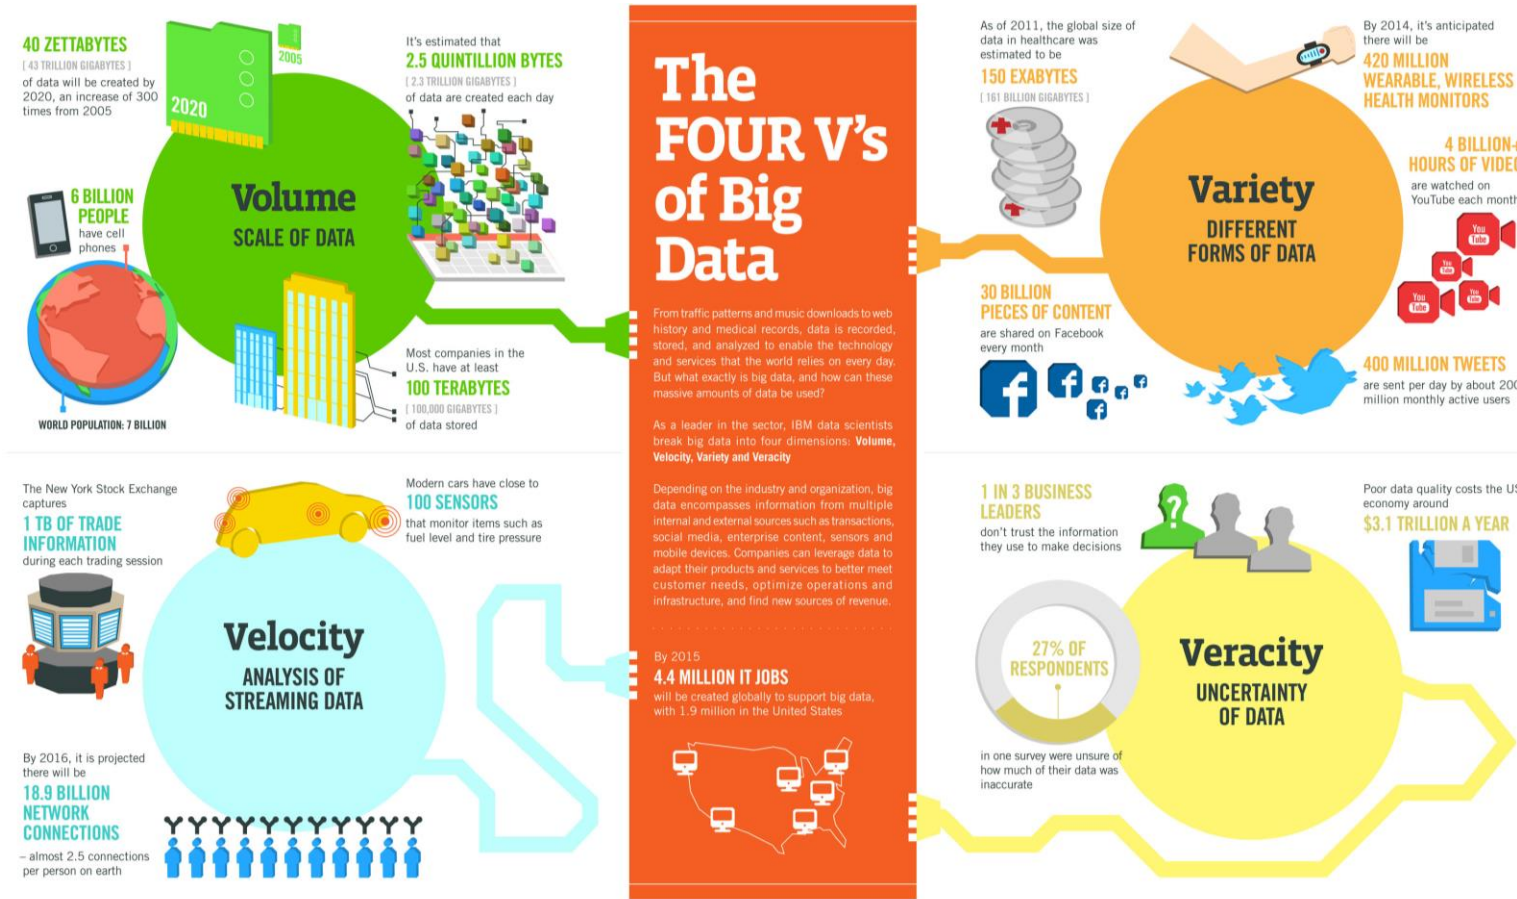
"Big data refers to data sets whose size is **beyond the ability of typical** database **software tools** to capture, store, manage and analyze." - *The McKinsey Global Institute, 2012*

"Big data is data sets that are so voluminous and complex that **traditional** data processing application **softwares are inadequate** to deal with them." - *Wikipedia*

When does data become "Big"?



The four "V's" of Big Data



http://www.ibmbigdatahub.com/sites/default/files/infographic_file/4-Vs-of-big-data.jpg

IBM

How do we process Big Data?

Main issues

- Where do we store the data?
- How do we process it?

Big Data greatly exceeds the size of the typical drives

- Even if a big drive existed, it would be too slow (at least for now)



Year: 1990
Size: 1.3 GB
Speed: 4,4 MB/s

5 minutes



Year: 2014
Size: 1 TB
Speed: 100 MB/s

3 hours



Year: 2015
Size: 1 TB
Speed: 600 MB/s

30 minutes

Scale up

Adding more processors and RAM, buying expensive and robust server

Pros

- Less power consumption than running multiple servers
- Cooling costs are less than scaling horizontally
- Generally less challenging to implement
- Less licensing costs
- Less networking equipment

Cons

- **PRICE**
- Greater risk of hardware failure causing bigger outages
- Generally severe vendor lock-in
- Not long-term: limited upgradeability in the future

Scale out

Adding more servers with less processors and RAM

Pros

- Much cheaper than scaling vertically
- New technologies simplify fault-tolerance and systems monitoring
- Easy to upgrade
- Usually cheaper
- Can literally scale infinitely

Cons

- More **licensing** fees
- Bigger footprint in the Data Center
- Higher utility cost (electricity and cooling)
- Possible need for more networking equipment (switches/routers)

Commodity hardware

You are not tied to expensive, proprietary offerings from a single vendor

You can choose standardized, commonly available hardware from a large range of vendors to build your cluster

Commodity \neq Low-end!

- Cheap components with high failure rate can be a false economy



The answer: cluster computing




100 hard disks? 2 mins to read 1TB

Cluster computing

Compute nodes are stored on racks

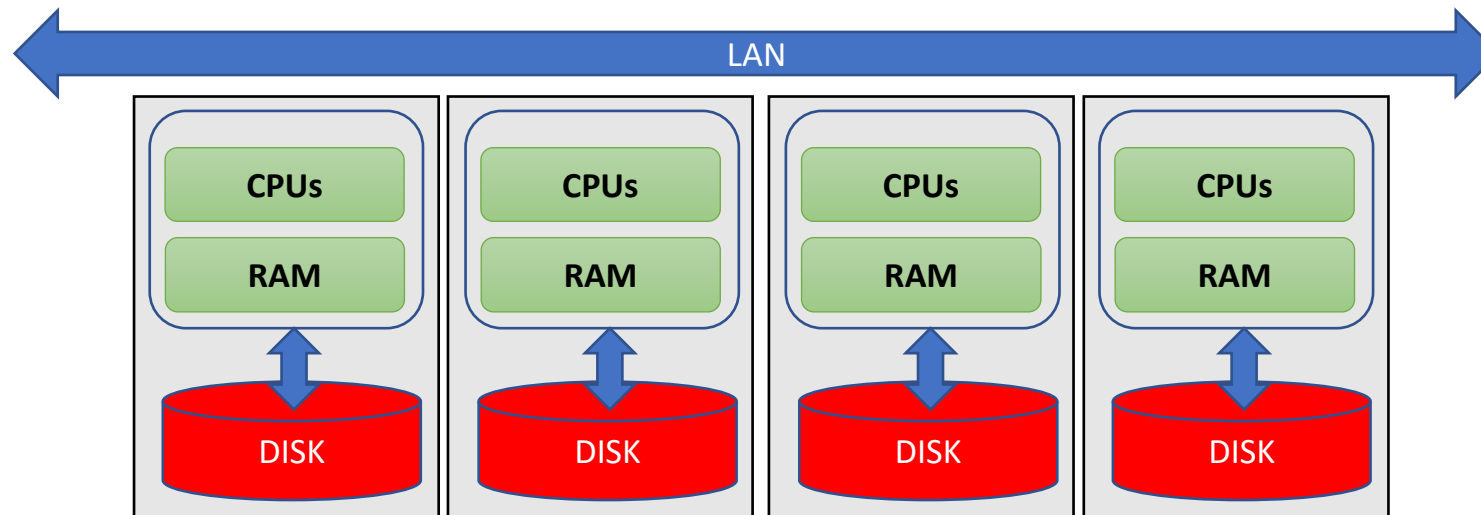
- 8–64 compute nodes on a rack
- There can be many racks of compute nodes
- The nodes on a single rack are connected by a network (typically gigabit Ethernet)
- Racks are connected by another level of network (or a switch)
 - The bandwidth of intra-rack communication is usually much greater than that of inter-rack communication



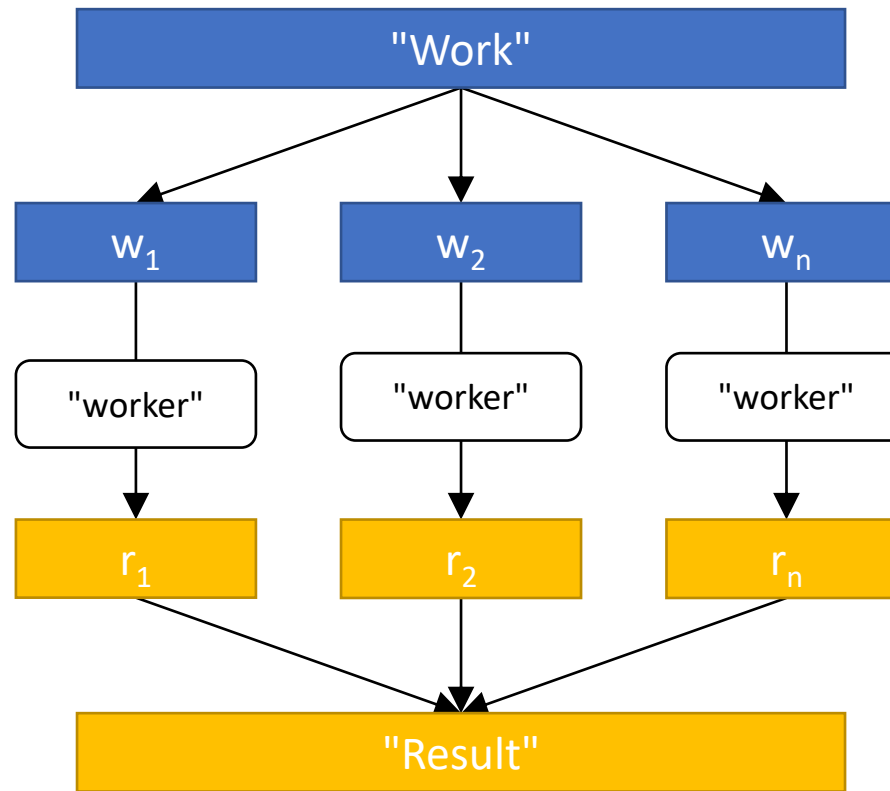
Cluster Computing Architecture

A computer cluster is a group of linked computers (nodes), working together closely so that in many respects they form a single computer

- Typically connected to each other through fast **LAN**
- **Every node is a system on its own**, capable of independent operations
 - Unlimited scalability, no vendor lock-in
- Number of nodes in the cluster \gg Number of CPUs in a node



Distributed computing: an old idea



Divide



Conquer

Parallelization Challenges

- How do we assign work units to workers?
- What if we have more work units than workers?
- What if workers need to share partial results?
- How do we aggregate partial results?
- How do we know all the workers have finished?
- What if workers die?




Risks?

Deadlock and starvation

Parallelization problems arise from:

- Communication between workers (e.g., to exchange state)
- Access to shared resources (e.g., data)

We need synchronization



But...

It is difficult to reason about parallelization

It is even more difficult to reason about concurrency

- At the scale of datacenters (even across datacenters)
- In the presence of failures
- In terms of multiple interacting services

Not to mention debugging...

The reality can be hard

- Lots of one-off solutions, custom code
- Write your own dedicated library, then program with it
- Burden on the programmer to explicitly manage everything

What is the solution?

Hide system-level details from the developers

- No more race conditions, lock contention, etc.
- No need to become hardcore techies

Separate the *what* from the *how*

- Developer specifies the computation that needs to be performed
- Execution framework (“runtime”) handles the actual execution

The datacenter **IS** the computer!




Disambiguation of MapReduce

"MapReduce is a programming model and an associated implementation for processing and generating large data sets. Users specify a map function that processes a key/value pair to generate a set of intermediate key/value pairs, and a reduce function that merges all intermediate values associated with the same intermediate key."

-- Dean J., Ghemawat S. (Google)

Hadoop MapReduce is an open-source implementation of the MapReduce programming model



Limitations of Map Reduce

Designed for batch processing

- Not suitable for iterative algorithms or interactive data mining

Strict paradigm


- Everything has to fit into Map and Reduce
- Complex algorithms will take multiple jobs and passes on hard disk

New hardware capabilities are not exploited

- Too much pressure on disk; RAM and multicore not adequately exploited

Too much complex

Limitations of Map Reduce



Spark

It is a **fast and general-purpose execution engine**

- **In-memory** data storage for very fast iterative queries
- Easy **interactive** data analysis
- Combines **different processing models** (machine learning, SQL, streaming, graph computation)
- Provides (not only) a MapReduce-like engine...
- ... but it's **up to 100x faster** than Hadoop MapReduce

Compatible with Hadoop's storage APIs

- Can run on top of a Hadoop cluster
- Can read/write to any database and any Hadoop-supported system, including HDFS, HBase, Parquet, etc.

What does Spark offer?

In-memory data caching

- HDD is scanned once, then data is written to/read from RAM

Lazy computations

- The job is optimized before its execution

Efficient pipelining

- Writing to HDD is avoided as much as possible

Spark pillars

Two main abstractions of Spark

RDD – Resilient Distributed Dataset

- An RDD is a collection of data items
- It is split into partitions
- It is stored in memory on the worker nodes of the cluster

DAG – Direct Acyclic Graph

- A DAG is a sequence of computations performed on data
- Each node is an RDD
- Each edge is a transformation of one RDD into another

RDD

RDDs are immutable distributed collection of objects

- **Resilient**: automatically rebuild on failure
- **Distributed**: the objects belonging to a given collection are split into *partitions* and spread across the nodes
 - RDDs can contain any type of Python, Java, or Scala objects
 - Distribution allows for scalability and locality-aware scheduling
 - Partitioning allows to control parallel processing

Fundamental characteristics (mostly from *pure functional programming*)

- **Immutable**: once created, it can't be modified
- **Lazily evaluated**: optimization before execution
- **Cacheable**: can persist in memory, spill to disk if necessary
- **Type inference**: data types are not declared but inferred (≠ dynamic typing)

RDD operations

RDDs offer two types of operations: *transformations* and *actions*

Transformations construct a new RDD from a previous one

- E.g.: map, flatMap, reduceByKey, filtering, etc.
- <https://spark.apache.org/docs/latest/programming-guide.html#transformations>

Actions compute a result that is either returned to the driver program or saved to an external storage system (e.g., HDFS)

- E.g.: saveAsTextFile, count, collect, etc.
- <https://spark.apache.org/docs/latest/programming-guide.html#actions>

RDD operations

RDDs are **lazily evaluated**, i.e., they are computed when they are used in an action

- Until no action is fired, the data to be processed is not even accessed

Example (in Python)

```
sc = new SparkContext
```

```
rddLines = sc.textFile("myFile.txt")
```

```
rddLines2 = rddLines.filter (lambda line: "some text" in line)
```

```
rddLines2.first()
```

} Transformations
} Action

There is no need to compute and store everything

- In the example, Spark simply scans the file until it finds the first matching line

DAG

Based on the user application and on the lineage graphs, Spark computes a **logical execution plan** in the form of a DAG

- Which is later transformed into a physical execution plan

The DAG (Directed Acyclic Graph) is **a sequence of computations performed on data**

- Nodes are **RDDs**
- Edges are operations on RDDs
- The graph is Directed: transformations from a partition A to a partition B
- The graph is Acyclic: transformations cannot return an old partition

Application decomposition

Application

- Single instance of SparkContext that stores data processing logic and schedules series of jobs, sequentially or in parallel

Job

- Complete set of transformations on RDD that finishes with action or data saving, triggered by the driver application


Stage

- Set of transformations that can be pipelined and executed by a single independent worker

Task

- Basic unit of scheduling: executes the stage on a single data partition


Application decomposition



DAG example

Word count in Scala


```
textFile = sc.textFile("hdfs://...")
```



DAG example

Word count in Scala


```
textFile = sc.textFile("hdfs://...")  
counts = textFile  
  .flatMap(line => line.split(" "))
```



DAG example

Word count in Scala


```
textFile = sc.textFile("hdfs://...")  
counts = textFile  
    .flatMap(line => line.split(" "))  
    .map(lambda word: (word, 1))
```



DAG example

Word count in Scala


```
textFile = sc.textFile("hdfs://...")  
counts = textFile  
    .flatMap(line => line.split(" "))  
    .map(lambda word: (word, 1))  
    .reduceByKey(lambda a, b: a + b)
```



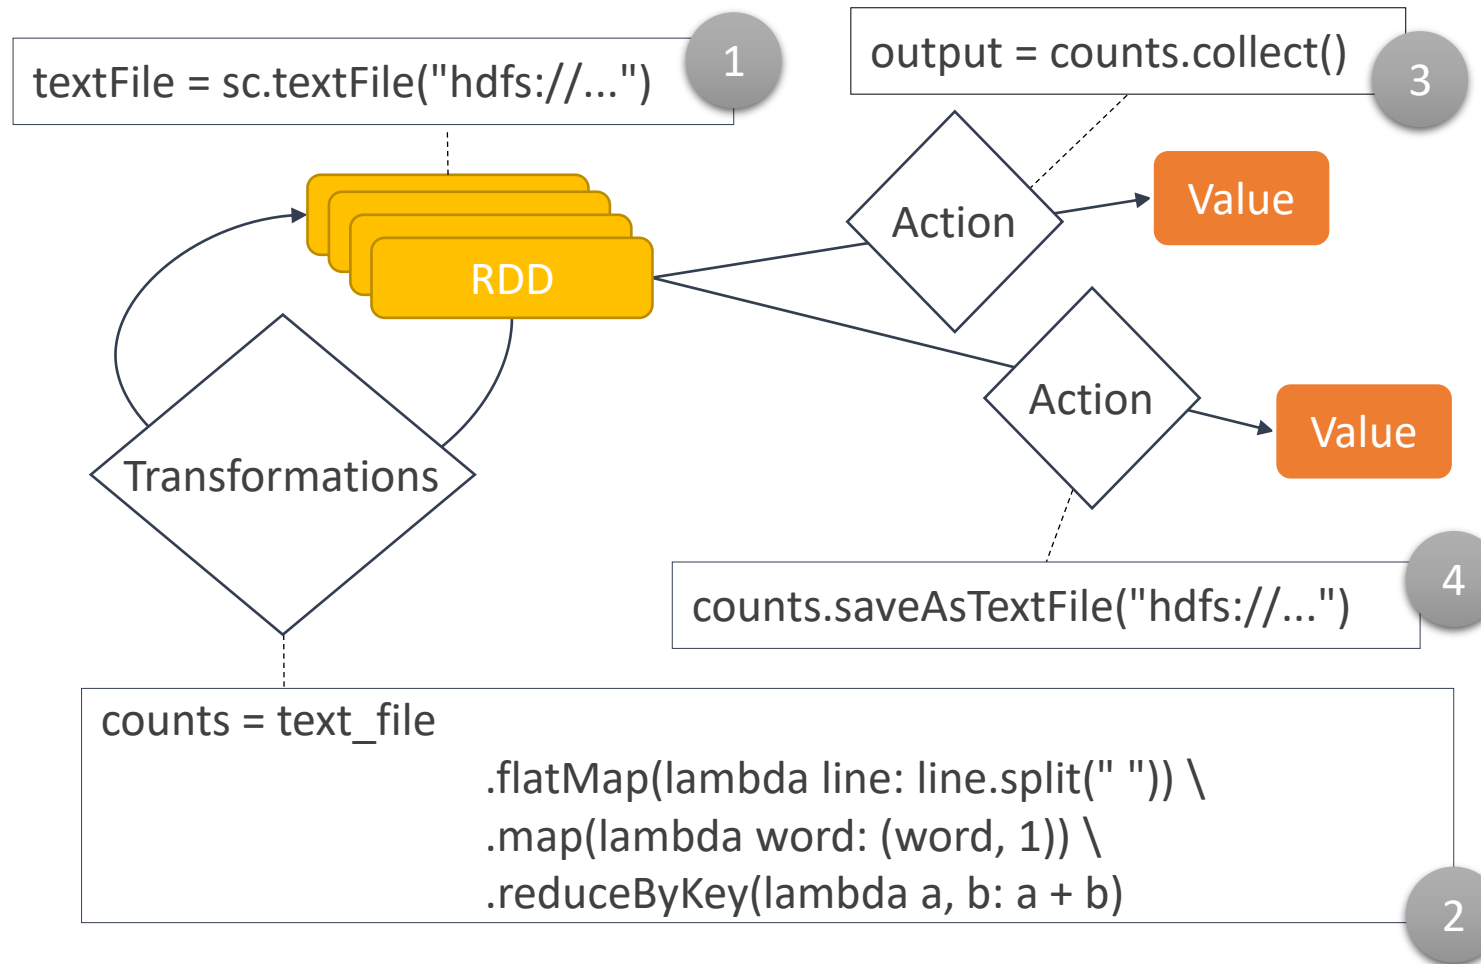
DAG example

Word count in Scala

```
textFile = sc.textFile("hdfs://...")  
counts = textFile  
    .flatMap(line => line.split(" "))  
    .map(lambda word: (word, 1))  
    .reduceByKey(lambda a, b: a + b)  
counts.saveAsTextFile("hdfs://...")
```



Conceptual representation



DataFrame and DataSet

RDDs are immutable distributed collection of objects

DataFrames and DataSets are immutable distributed collection of records organized into named columns (i.e., a table)

- **Simply put, RDDs with a schema attached**
- Support both relational and procedural processing (e.g., SQL, Scala)
- Support complex data types (struct, array, etc.) and user defined types
- Cached using columnar storage

Can be built from many different sources

- DBMSs, files, other tools (e.g., Hive), RDDs

Type conformity is checked

- *At compile time* for DataSets; *at runtime* for DataFrames

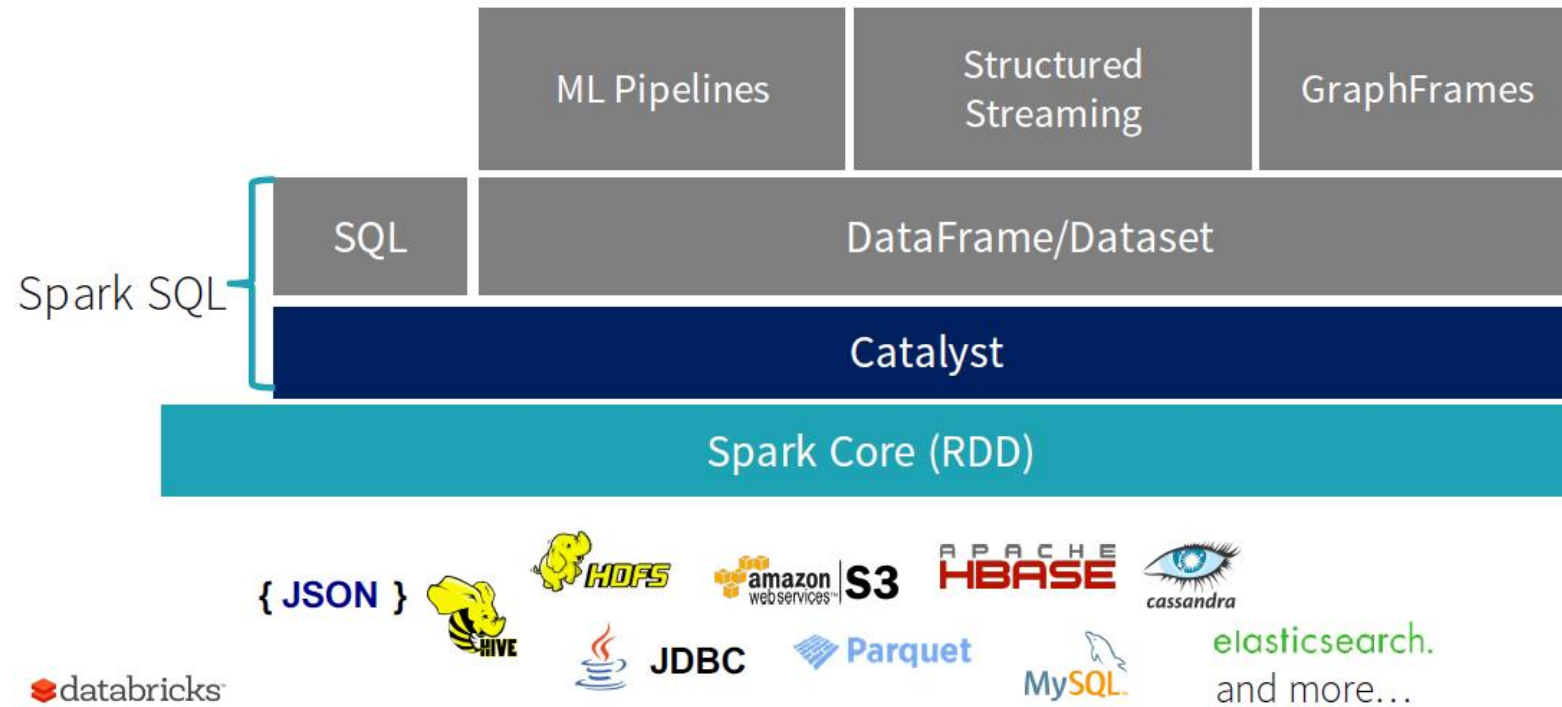
DataFrame and DataSet

Still lazily evaluated...

...but supports under-the-hood optimizations and code generation

- **Catalyst optimizer creates optimized execution plans**
 - IO optimizations such as skipping blocks in parquet files
 - Logic push-down of selection predicates
- JVM code generation for all supported languages
 - Even non-native JVM languages; e.g., Python

Spark structured



Why structure?


Cons

- **Structure imposes some limits**
 - RDDs enable any computation through user defined functions

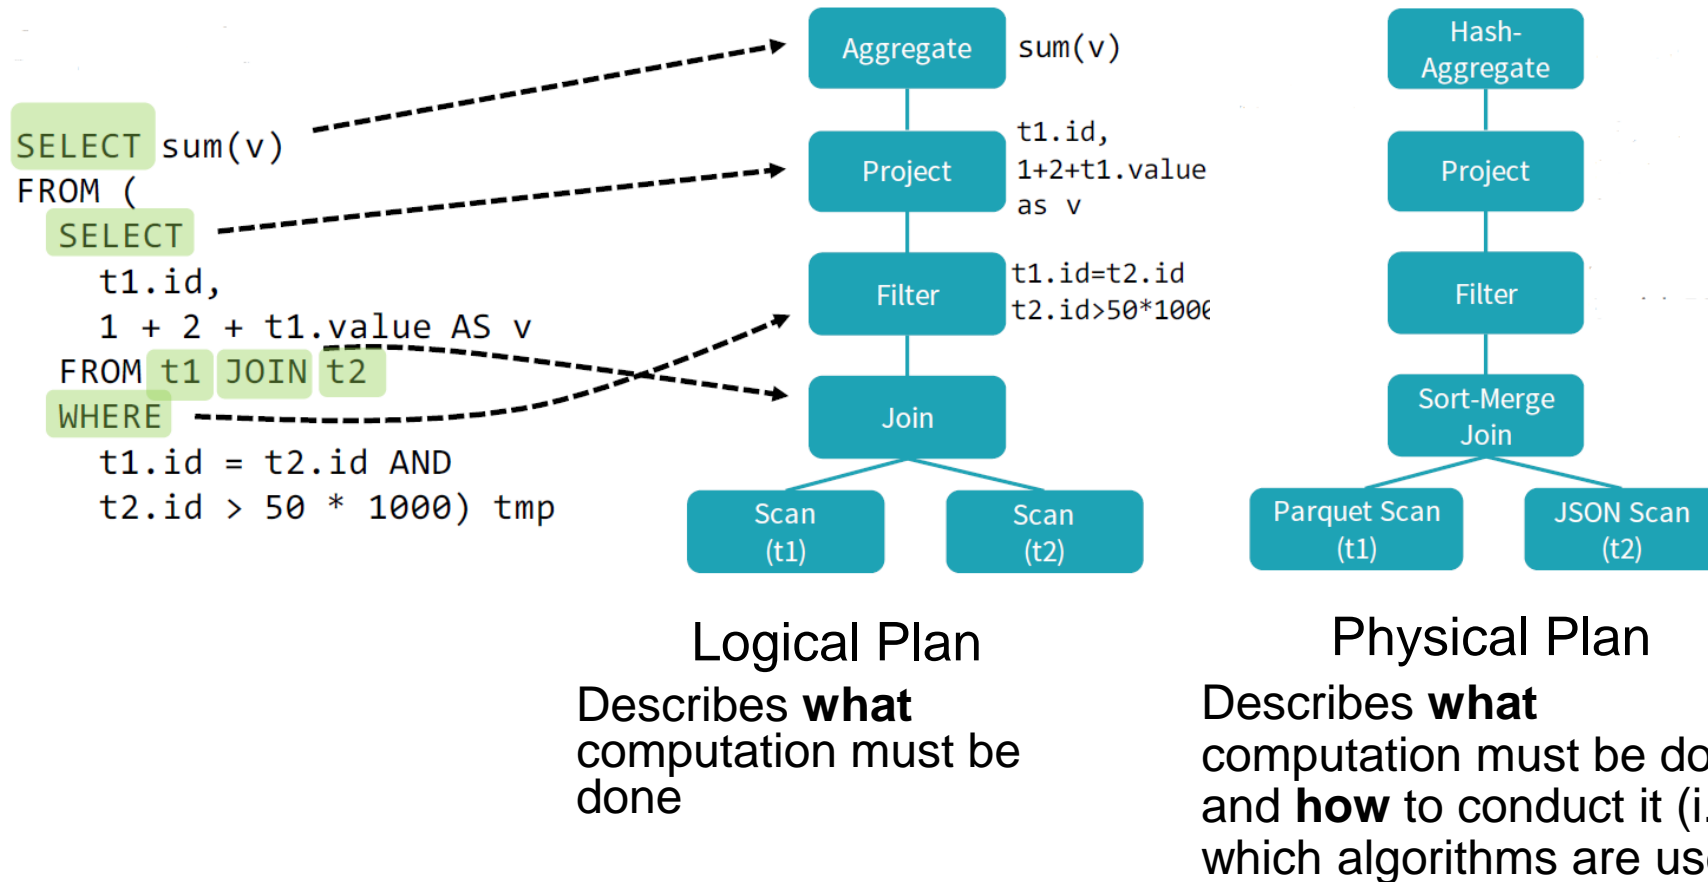
Pros

- The most common computations are supported
- Language simplicity
- **Opens the room to optimizations**
 - Hard to optimize a user defined function

Catalyst



Logical and Physical Plan



Logical optimization

Based on rules

- **A rule is a function** that can be applied on a portion of the logical plan

Implemented as Scala functions

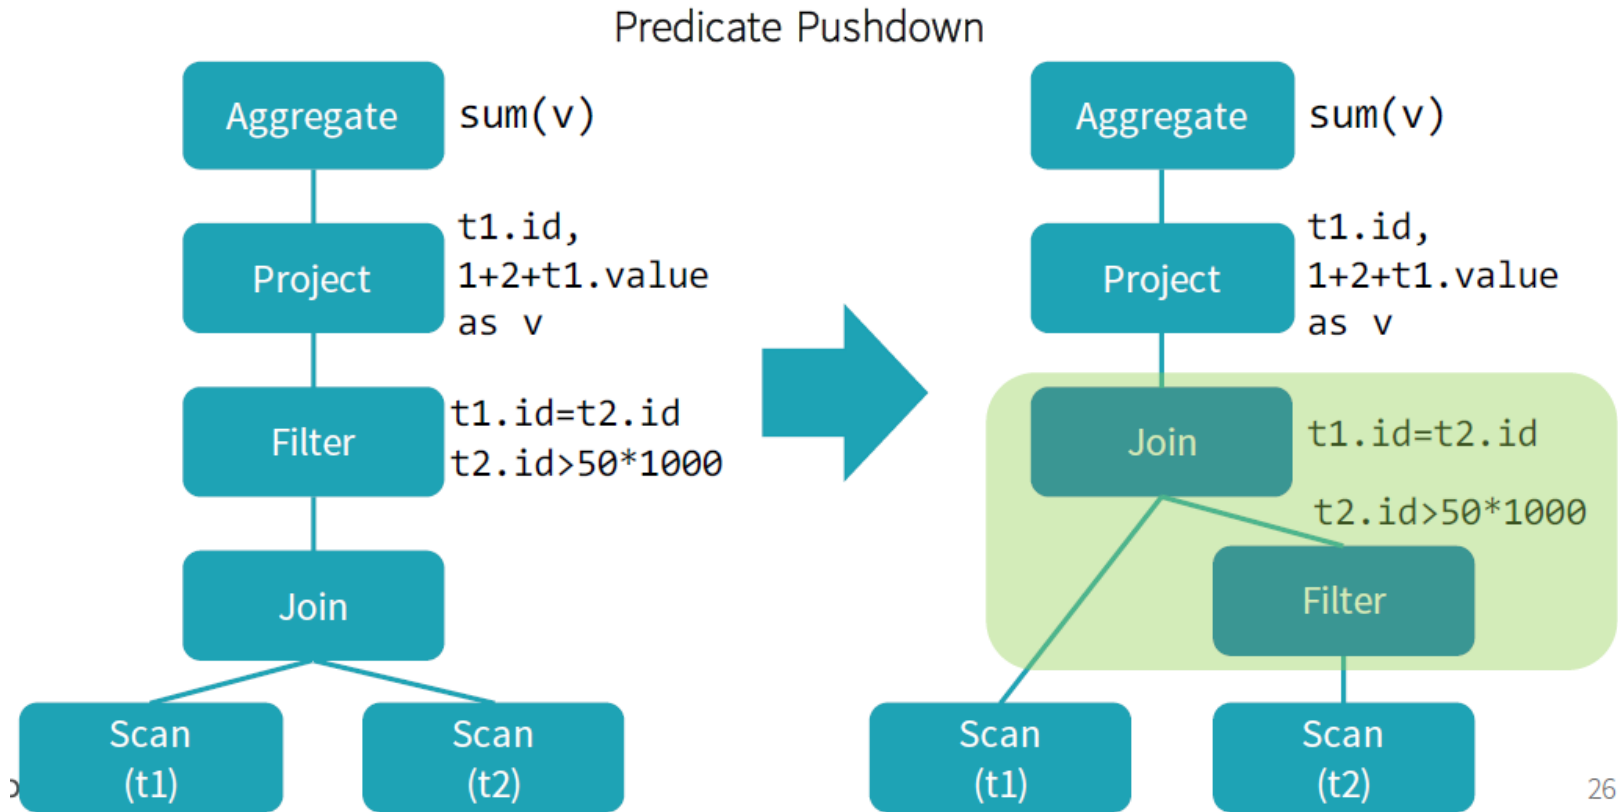
```
val expression: Expression = ...
expression.transform {
  case Add(Literal(x, IntegerType), Literal(y, IntegerType)) =>
    Literal(x + y)
}
```

Several types of rules

- **Constant folding**: resolve constant expressions at compile time
- **Predicate pushdown**: push selection predicates close to the sources
- **Column pruning**: project only the required column
- **Join reordering**: change the order of join operations

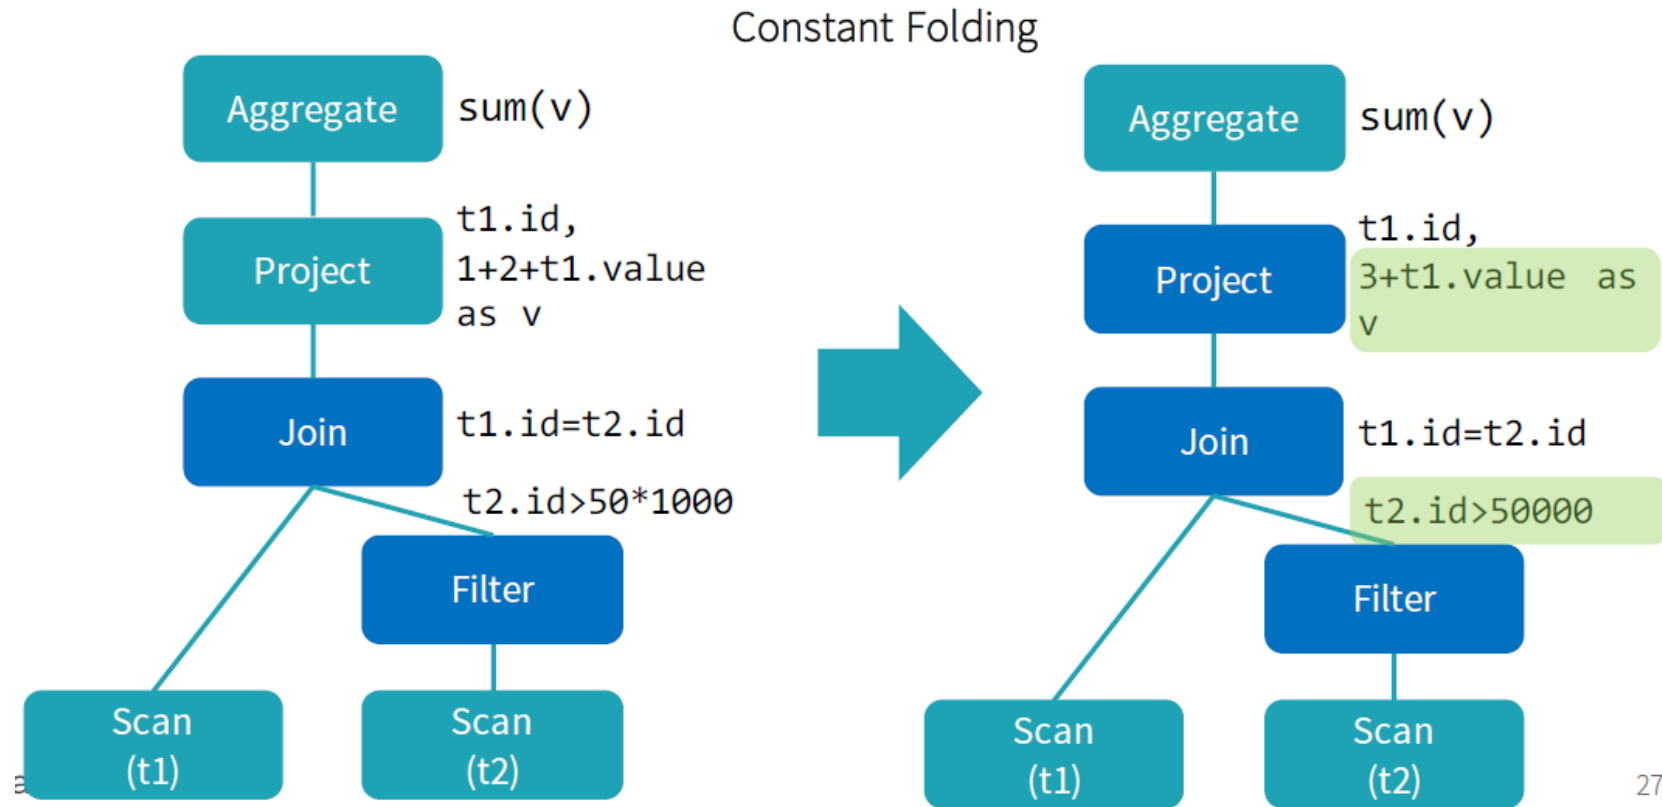
Applied recursively and iteratively until the plan reaches a *fixed point*

Logical optimization



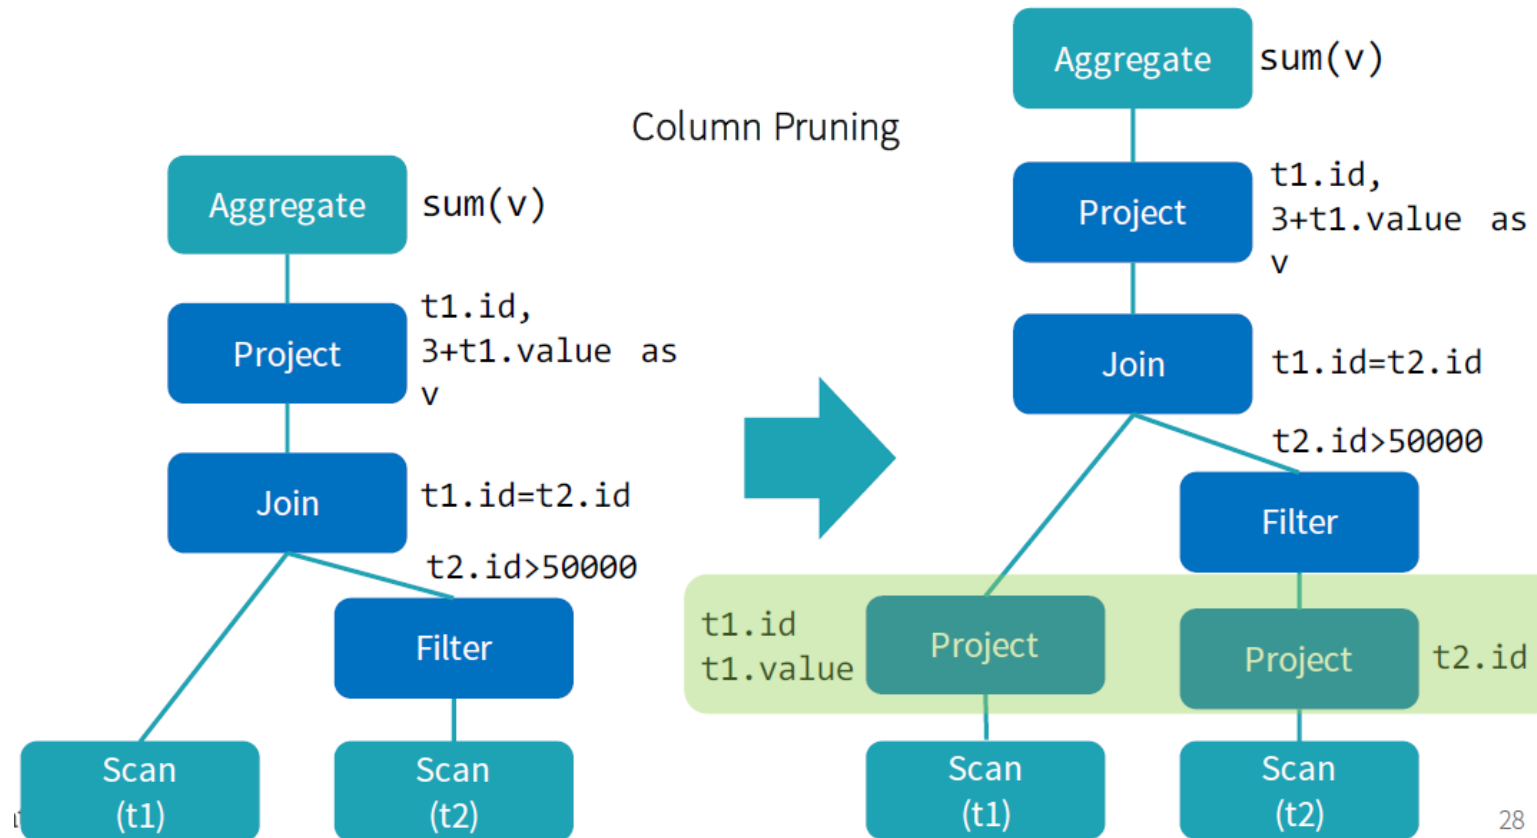
26

Logical optimization



27

Logical optimization



28

What determines the cost?

Catalyst only considers the size and the cardinality of tables

Other important factors

- Network throughput
- Disk throughput
- Allocation of resources
 - Number of Executors
 - Number of Cores per Executor
- Allocation of tasks
 - Data locality probability

Research work: defining a more accurate, probabilistic cost model


- Baldacci, L., & Golfarelli, M. (2018). A cost model for Spark SQL. *IEEE Transactions on Knowledge and Data Engineering*, 31(5), 819-832.
- Gallinucci, E., & Golfarelli, M. (2019). SparkTune: tuning Spark SQL through query cost modeling. *EDBT 2019*: 546-549.

Spark architecture

Spark uses a *master/slave architecture* with one central coordinator (*driver*) and many distributed workers (*executors*)

- The driver and each executor are independent Java processes
- Together they form a Spark *application*

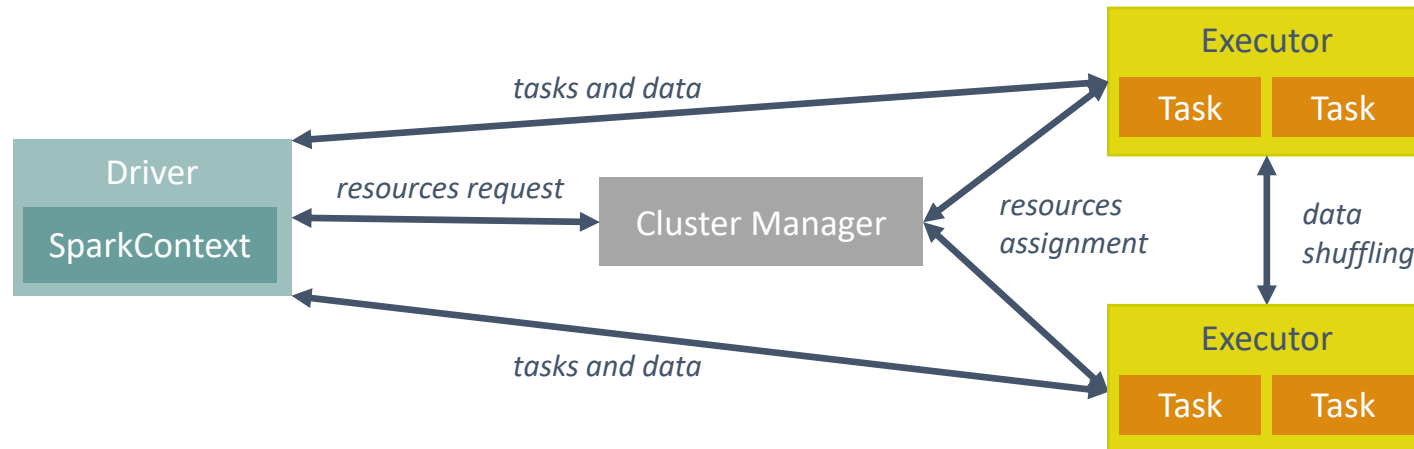
The architecture is independent of the cluster manager that Spark runs on



Spark architecture

Executor: a process responsible for executing the received tasks

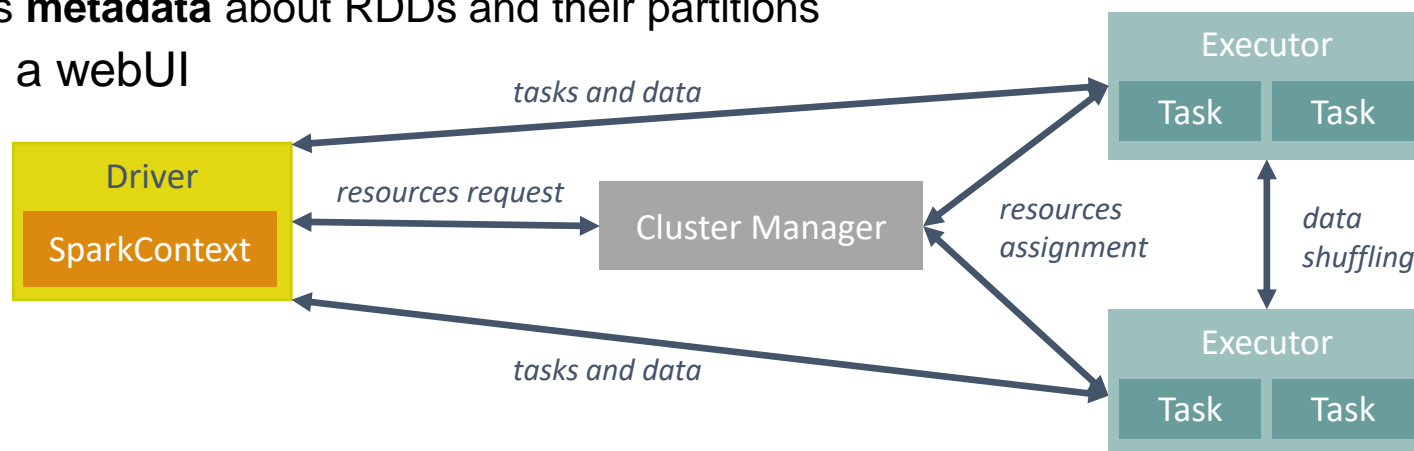
- Each spark application can have (and usually has) multiple executors, and each worker node can host many executors
- Typically runs for the entire duration of the application
- Stores (caches) RDD data in JVM heap
- **Tasks** are the smallest unit of work and are carried out by executors



Spark architecture

Driver Program (a.k.a. *Spark Driver*, or simply *Driver*)

- Each spark application can only have one driver (entry point of Spark Shell)
- Converts user program into tasks
 - Creates the **SparkContext**, i.e., the object that handles communications
 - Computes the logical **DAG** of operations and converts it into a physical **execution plan**
- Schedules tasks on executors
 - Has a **complete view** of the available executors and schedules tasks on them
 - Stores **metadata** about RDDs and their partitions
- Launches a webUI



In action!



Enter the folder `03-BigData`



Double click on `run.bat`



Open the browser




Copy and paste the link 127.0.0.1:8080



Enter the notebook `04-BigData`



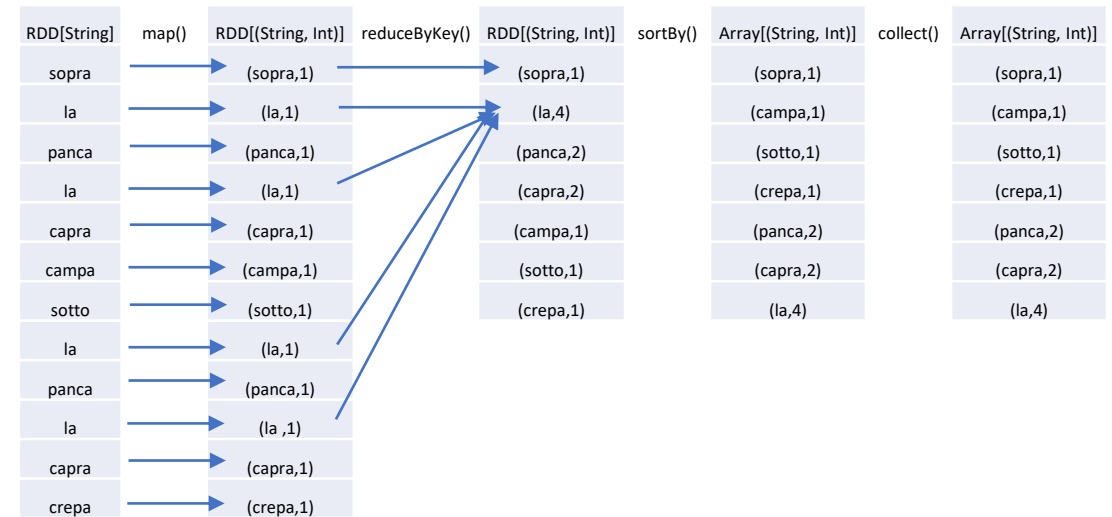
2



Word count

rdd

```
.map(s => (s, 1))  
.reduceByKey((a, b) => a + b)  
.sortBy(x => x._2) tuples  
.collect()
```



Foodmart

foodmartDataset

```
.groupBy("category")  
.sum("unit sales")  
.orderBy("sum(unit sales)")  
.show()
```

foodmartDataset			groupBy(category)	foodmartDataset			sum(unit_sales)	foodmartDataset		orderBy(sum(unit_sales))	foodmartDataset		sum(unit_s
product	category	unit_sales		product	category	unit_sales		category	sum(unit_sales)		category	sum(unit_s	ales)
beer	drink	2		bread	food	4		food	9		drink	5	
cola	drink	3		pizza	food	5		drink	5		food	9	
bread	food	4		beer	drink	2							
pizza	food	5		cola	drink	3							

Spark

Suggested reading and resources

